

DD2418 Language Engineering

5a: Text classification

Johan Boye, KTH

March 30, 2020

Is this spam?

Hello

Want to gain good money with no effort and spend little time on that?

As of now, you can gain good money with our binary options service. Remember that for this month we have paid \$11,000,000 to our users

For your awesome, we have prepared a bonus for you: \$1,000.

Everything is ready. Just click SIGN UP and start earning!

SIGN UP

© 2018 All Rights Reserved

Is this spam?

Hello

Want to gain good money with no effort and spend little time on that?

As of now, you can get good money with our binary options service. Remember that for this month we have paid **\$11,000,000** to our users

For your awesome, we have prepared a bonus for you **\$1,000**.

Everything is ready. Just click SIGN UP and start earning!

SIGN UP

© 2018 All Rights Reserved

Is this spam?

From: Mihaela Ioana <contact@icef4.info>

Sent: den 23 september 2018 22:11

To: info@icef4.org

Subject: You are invited to Amsterdam Science Park on 14 December

Dear colleagues

You are now welcome to participate to 4th International Conference on Engineering in the premises of **Amsterdam University Science Park**, with your abstract, full paper and poster, as in-person or virtual.

The conference topics include **Engineering** fields in Aerospace, Biomedical, Chemical, Civil, Electrical, Computer Science, Environmental, Industrial, Materials, Mechanical, Nuclear as well as Formal Sciences such as Mathematics and Physics.

The accepted full papers will be published in the proceedings book with ISBN and in indexed journals by **De Gruyter**.

You can register with your **abstract** from now until 25 November 2018. The abstract book with ISBN will be published before the conference. Sending paper is optional. Full Proceedings and Journal will be published within 2 months after the conference.

The **fee starts from €170** and it varies depending on your preferences such as participation type, journal publishing etc.

You will receive an acceptance letter with an invoice within 7 days after submission based on your abstract. You will receive the receipt and an invitation letter upon the payment. The committee makes sure to accept your optional full paper afterwards. This is helpful for you to make early travel planning.

For more details and registration, please visit <http://icef4.org>

Looking forward to meeting you at Amsterdam Science Park.

Best Regards
Mihaela Ioana,
Secretary ICEF IV

Is this spam?

From: Mihaela Ioana <contact@icef4.info>

Sent: den 23 september 2018 22:11

To: info@icef4.org

Subject: You are invited to Amsterdam Science Park on 14 December

Dear colleagues

You are now welcome to participate to 4th International Conference on Engineering in the premises of **Amsterdam University Science Park**, with your abstract, full paper and poster, as in-person or virtual.

The conference topics include **Engineering** fields in Aerospace, Biomedical, Chemical, Civil, Electrical, Computer Science, Environmental, Industrial, Materials, Mechanical, Nuclear as well as Formal Sciences such as Mathematics and Physics.

The accepted full papers will be published in the proceedings book with ISBN and in indexed journals by **De Gruyter**.

You can register with your **abstract** from now until 25 November 2018. The abstract book with ISBN will be published before the conference. Sending paper is optional. Full Proceedings and Journal will be published within 2 months after the conference.

The fee starts from €170 and it varies depending on your preferences such as participation type, journal publishing etc.

You will receive an acceptance letter with an invoice within 7 days after submission based on your abstract. You will receive the receipt and an invitation letter upon the payment. The committee makes sure to accept your optional full paper afterwards. This is helpful for you to make early travel planning.

For more details and registration, please visit <http://icef4.org>

Looking forward to meeting you at Amsterdam Science Park.


Best Regards
Mihaela Ioana,
Secretary ICEF IV

Are these reviews positive or negative?

Uncanny, haunting, I must have read this novel at the right time for me as it found a sure spot under my skin and disturbed my normally peaceful sleep. positive review

If you are reading this because it is your darling fragrance, please wear it at home exclusively, and tape the windows shut. humouristic negative

Who wrote Henry VI, part 3?




Shakespeare or Marlowe? In 1994, two researchers using neural network techniques found strong support that Marlowe wrote the original.

Merriam, T. V., & Matthews, R. A. (1994). Neural computation in stylometry II: An application to the works of Shakespeare and Marlowe. *Literary and Linguistic Computing*, 9(1), 1-6.

By **MAYA KAUFMAN** | CBS NEWS | August 6, 2015, 6:00 AM

Language detectives make the web less anonymous



/ ISTOCKPHOTO

Text classification

We want to automatically assign a *label* to a document (=book, article, web page, paragraph, sentence).

It is assumed that every document has exactly one label.

A machine learning formulation

Given:

- a set of documents $\{d_1, d_2, \dots, d_m\}$
- a set of labels (classes) $\{y_1, y_2, \dots, y_n\}$
- a training set of labeled documents $\{(d_1, y_{d_1}), \dots, (d_m, y_{d_m})\}$

produce:

- a classifier $h(d) = y$

How should we represent d ?

Bag-of-words representation

A vector of word counts is called a *bag of words*.

e.g 2 occurrences of "man"
3 occurrences of "dog" $(2, 3, 1, 0, 0, 0, 1, \dots, 0, 1)$

The vector has length n = the number of unique words in the training set.

Note that word order is not represented!

Bag-of-words representation

You have the following corpus. Create word vectors for all documents!

to be is to do
to do is to be
do be do be do
to be or not to be
I do I do I do I do I do

Sartre

Kant

Sinatra

Shakespeare

ABBA

Bag-of-words representation


You have the following corpus. Create word vectors for all documents!

					to	be	is	to	do	
					to	do	is	to	be	
					do	be	do	be	do	
					to	be	or	not	to	be
I	do	I	do	I	do	I	do	I	do	

to	be	is	do	or	not	I							
(2	,	1	,	1	,	1	,	0	,	0	,	0)
(2	,	1	,	1	,	1	,	0	,	0	,	0)
(0	,	2	,	0	,	3	,	0	,	0	,	0)
(2	,	2	,	0	,	0	,	1	,	1	,	0)
(0	,	0	,	0	,	5	,	0	,	0	,	5)

Comparing documents

Through the bag-of-words representation, documents can be seen as *points* or *vectors* in a multi-dimensional space.



Comparing vectors

We can compare our documents vectors in different ways.

Euclidean distance:

$$\|(x - y)\|_2 = \sqrt{\sum_{i=1}^n (x_i - y_i)^2}$$

Manhattan distance :

$$\|(x - y)\|_1 = \sum_{i=1}^n |x_i - y_i|$$

Cosine similarity :

can see it as a measure of differences in angle.

Smaller angle -> More similar

$$\cos(x, y) = \frac{x \cdot y}{\|x\| \|y\|}$$

more fruitful to compare similarity than differences

will get a number between 0 and 1.

Comparing vectors


refer to slide 12

$$\begin{aligned}\cos(\text{Sartre}, \text{Kant}) &= \\ \cos((2, 1, 1, 1, 0, 0, 0), (2, 1, 1, 1, 0, 0, 0)) &= \\ \frac{2 \cdot 2 + 1 \cdot 1 + 1 \cdot 1 + 1 \cdot 1 + 0 \cdot 0 + 0 \cdot 0 + 0 \cdot 0}{\sqrt{7} \sqrt{7}} &= \\ 1\end{aligned}$$

$$\begin{aligned}\cos(\text{Sartre}, \text{Sinatra}) &= \\ \cos((2, 1, 1, 1, 0, 0, 0), (0, 2, 0, 3, 0, 0, 0)) &= \\ \frac{2 \cdot 0 + 1 \cdot 2 + 1 \cdot 0 + 1 \cdot 3 + 0 \cdot 0 + 0 \cdot 0 + 0 \cdot 0}{\sqrt{7} \sqrt{13}} &= \\ 0.52\end{aligned}$$

$$\begin{aligned}\cos(\text{Shakespeare}, \text{ABBA}) &= \\ \cos((2, 2, 0, 0, 1, 1, 0), (0, 0, 0, 5, 0, 0, 5)) &= \\ \frac{2 \cdot 0 + 2 \cdot 0 + 0 \cdot 0 + 0 \cdot 5 + 1 \cdot 0 + 1 \cdot 0 + 0 \cdot 5}{\sqrt{10} \sqrt{50}} &= \\ 0\end{aligned}$$

k -nearest neighbour



Using cosine similarity, we have:

- 1-nearest neighbour $\Rightarrow ? = \text{ham}$
- 3-nearest neighbour $\Rightarrow ? = \text{spam}$
- 5-nearest neighbour $\Rightarrow ? = \text{ham}$

Dimensionality reduction

is useful to remove the curse of dimensionality

Do we have to consider all words?

No! Usually it is beneficial to remove so-called very common words (*stop words*).
doesn't add very much to semantics

These can be defined *grammatically* (pronouns, prepositions, determiners, ...), or *statistically* (words that are common in the corpus).

To Sherlock Holmes she is always *the* woman. I have seldom heard him mention her under any other name. In his eyes she eclipses and predominates the whole of her sex. It was not that he felt any emotion akin to love for Irene Adler. All emotions, and that one particularly, were abhorrent to his cold, precise but admirably balanced mind. He was, I take it, the most perfect reasoning and observing machine that the world has seen, but as a lover he would have placed himself in a false position. He never spoke of the softer passions, save with a gibe and a sneer. They were admirable things for the observer—excellent for drawing the veil from men's motives and actions. But for the trained reasoner to admit such intrusions into his own delicate and finely adjusted temperament was to introduce a distracting factor which might throw a doubt upon all his mental results. Grit in a sensitive instrument, or a crack in one of his own high-power lenses, would not be more disturbing than a strong emotion in a nature such as his. And yet there was but one woman to him, and that woman was the late Irene Adler, of dubious and questionable memory.

To Sherlock Holmes she is always *the* woman. I have seldom heard him mention her under any other name. In his eyes she eclipses and predominates the whole of her sex. It was not that he felt any emotion akin to love for Irene Adler. All emotions, and that one particularly, were abhorrent to his cold, precise but admirably balanced mind. He was, I take it, the most perfect reasoning and observing machine that the world has seen, but as a lover he would have placed himself in a false position. He never spoke of the softer passions, save with a gibe and a sneer. They were admirable things for the observer—excellent for drawing the veil from men's motives and actions. But for the trained reasoner to admit such intrusions into his own delicate and finely adjusted temperament was to introduce a distracting factor which might throw a doubt upon all his mental results. Grit in a sensitive instrument, or a crack in one of his own high-power lenses, would not be more disturbing than a strong emotion in a nature such as his. And yet there was but one woman to him, and that woman was the late Irene Adler, of dubious and questionable memory.


10 a
9 the
9 and
6 to
6 his
5 was
5 in
4 that
4 of
4 he
4 for
4 but
3 woman
3 one
2 would
2 were
2 such
2 she
2 own
2 not


To Sherlock Holmes she is always *the* woman. I have seldom heard him mention her under any other name. In his eyes she eclipses and predominates the whole of her sex. It was not that he felt any emotion akin to love for Irene Adler. All emotions, and that one particularly, were abhorrent to his cold, precise but admirably balanced mind. He was, I take it, the most perfect reasoning and observing machine that the world has seen, but as a lover he would have placed himself in a false position. He never spoke of the softer passions, save with a gibe and a sneer. They were admirable things for the observer—excellent for drawing the veil from men's motives and actions. But for the trained reasoner to admit such intrusions into his own delicate and finely adjusted temperament was to introduce a distracting factor which might throw a doubt upon all his mental results. Grit in a sensitive instrument, or a crack in one of his own high-power lenses, would not be more disturbing than a strong emotion in a nature such as his. And yet there was but one woman to him, and that woman was the late Irene Adler, of dubious and questionable memory.

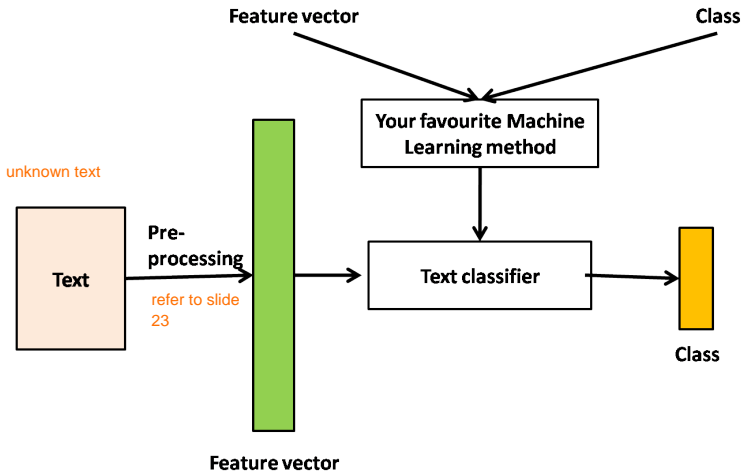

3 woman
2 own
2 irene
2 emotion
2 adler
1 yet
1 world
1 whole
1 veil
1 trained
1 throw
1 things
1 the
1 temperament
1 take
1 such
1 strong
1 spoke
1 softer
1 sneer

To Sherlock Holmes she is always *the* woman. I have seldom **hear** him mention her under any other name. In his **eye** she **eclipse** and **predominate** the whole of her sex. It was not that he **feel** any emotion akin to love for Irene Adler. All **emotion** , and that one particularly, were abhorrent to his cold, precise but admirably balanced mind. He was, I take it, the most perfect reasoning and observing machine that the world has **see** , but as a lover he would have **place** himself in a false position. He never **speak** of the **soft passion** , save with a gibe and a sneer. They were admirable things for the observer—excellent for **draw** the veil from **man motive** and **action** . But for the **train** reasoner to admit such **intrusion** into his own delicate and finely **adjust** temperament was to introduce a distracting factor which might throw a doubt upon all his mental **result** . Grit in a sensitive instrument, or a crack in one of his own high-power **lens** , would not be more disturbing than a strong emotion in a nature such as his. And yet there was but one woman to him, and that woman was the late Irene Adler, of dubious and questionable memory.

3 woman
2 own
2 irene
2 emotion
2 adler
1 yet
1 world
1 whole
1 veil
1 train
1 throw
1 things
1 the
1 temperament
1 take
1 such
1 strong
1 speak
1 soft
1 sneer







More dimensionality reduction

- Remove stop words
- Remove rare words
- Case folding turn everything into lower case
- Remove diacritics
- Spell checking is a dimensionality reduction technique. Misspelling creates a lot of version and hence vectors of word
- Lemmatization
- Replace numbers with `<number>` special tag
- Replace mail addresses with `<mail>`, web addresses with `<web>`, etc.
- Normalize punctuation (e.g. “ becomes ”)

There are also more advanced dimensionality reduction methods. More about them later in the course.

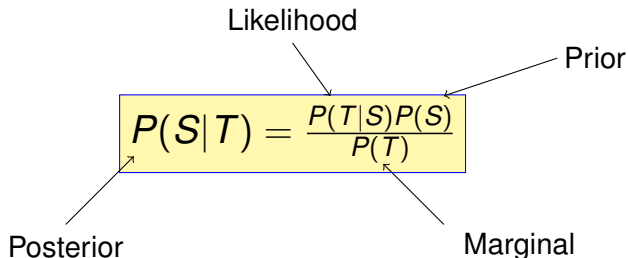
DD2418 Language Engineering

5b: Naive Bayes text classification

Johan Boye, KTH

Bayes' theorem

Bayes' theorem expresses how beliefs should be updated in the light of new evidence.



A diagram illustrating Bayes' theorem. The equation $P(S|T) = \frac{P(T|S)P(S)}{P(T)}$ is displayed within a yellow rectangular box. Four arrows point from labels to parts of the equation: 'Likelihood' points to $P(T|S)$, 'Prior' points to $P(S)$, 'Posterior' points to $P(S|T)$, and 'Marginal' points to $P(T)$.

$$P(S|T) = \frac{P(T|S)P(S)}{P(T)}$$

Labels and arrows:

- Likelihood points to $P(T|S)$
- Prior points to $P(S)$
- Posterior points to $P(S|T)$
- Marginal points to $P(T)$

Naive Bayes classifier

Given a document with words $w_1 \dots w_n$, what is the most likely label y ?

$$\arg \max_y P(y|w_1 \dots w_n)$$

want to find the class label y that maximize the probability

Rewrite and simplify:

$$\arg \max_y P(y|w_1 \dots w_n) \stackrel{\text{Bayes' theorem}}{=}$$

$$\arg \max_y \frac{P(w_1 \dots w_n|y)P(y)}{P(w_1 \dots w_n)} =$$

$$\arg \max_y P(w_1 \dots w_n|y)P(y) \stackrel{\text{conditional independence assumption}}{=}$$

$$\arg \max_y P(w_1|y) \dots P(w_n|y)P(y)$$

Parameter estimation

Maximum likelihood estimation:

$$P(y) = \frac{\text{\# of documents of label } y}{\text{total number of documents}}$$

$$P(w_i|y) = \frac{\text{\# of times } w_i \text{ appears in } y}{\text{\# of tokens in all } y\text{-documents}}$$

Example

2 labels: +, -.

4 documents: A, B, C, D .

	cold	cough	fever	flu	influenza	pneumonia	temperature	
A	0	1	1	1	0	0	1	+
B	0	1	0	1	0	1	2	-
C	0	1	0	2	0	0	1	+
D	1	0	0	1	0	0	0	+

Construct a Naive Bayes classifier, and use it to classify the document “*temperature*”.

Example

2 labels: +, -.

4 documents: A, B, C, D .

	cold	cough	fever	flu	influenza	pneumonia	temperature	
A	0	1	1	1	0	0	1	+
B	0	1	0	1	0	1	2	-
C	0	1	0	2	0	0	1	+
D	1	0	0	1	0	0	0	+

$$P(+) = \frac{3}{4}, \quad P(-) = \frac{1}{4}, \quad P(\text{temperature}|+) = \frac{2}{10}, \quad P(\text{temperature}|-) = \frac{2}{5}$$

Predicted label

$$P(+|\text{temperature}) = \frac{1}{Z} P(+) P(\text{temperature}|+) = \frac{1}{Z} \cdot \frac{3}{4} \cdot \frac{2}{10} = \frac{1}{Z} \cdot \underline{0.15}$$

$$P(-|\text{temperature}) = \frac{1}{Z} P(-) P(\text{temperature}|-) = \frac{1}{Z} \cdot \frac{1}{4} \cdot \frac{2}{5} = \frac{1}{Z} \cdot \underline{0.1}$$

where $Z = P(\text{temperature})$, which can be ignored (why?).

The predicted label is the one with the highest probability:

$$\arg \max_{y \in \{+, -\}} P(y|\text{temperature}) = +$$

Another example

2 labels: +, -.

4 documents: *A*, *B*, *C*, *D*.

	cold	cough	fever	flu	influenza	pneumonia	temperature	
<i>A</i>	0	1	1	1	0	0	1	+
<i>B</i>	0	1	0	1	0	1	2	-
<i>C</i>	0	1	0	2	0	0	1	+
<i>D</i>	1	0	0	1	0	0	0	+

Construct a Naive Bayes classifier, and use it to classify the document “*flu influenza*”.

Another example, cont.

	cold	cough	fever	flu	influenza	pneumonia	temperature	
<i>A</i>	0	1	1	1	0	0	1	+
<i>B</i>	0	1	0	1	0	1	2	-
<i>C</i>	0	1	0	2	0	0	1	+
<i>D</i>	1	0	0	1	0	0	0	+

$$P(+) = \frac{3}{4}, \quad P(-) = \frac{1}{4}, \quad P(\text{flu}|+) = \frac{4}{10}, \quad P(\text{flu}|-) = \frac{1}{5}$$

However:

$$P(\text{influenza}|+) = P(\text{influenza}|-) = 0$$

Laplace smoothing

	cold	cough	fever	flu	influenza	pneumonia	temperature	
A	0	1	1	1	0	0	1	+
B	0	1	0	1	0	1	2	-
C	0	1	0	2	0	0	1	+
D	1	0	0	1	0	0	0	+

Laplace smoothing: Add 1 occurrence of every word to every label

$$P(+)P(\text{flu}|+)P(\text{influenza}|+) = \frac{3}{4} \cdot \frac{4+1}{10+7} \cdot \frac{1}{10+7} = \underline{0.013}$$

$$P(-)P(\text{flu}|-)P(\text{influenza}|-) = \frac{1}{4} \cdot \frac{1+1}{5+7} \cdot \frac{1}{5+7} = \underline{0.003}$$

Predicted label

$$P(+|\text{flu influenza}) = \frac{1}{Z} P(+) P(\text{flu}|+) P(\text{influenza}|+) = \frac{1}{Z} \cdot \frac{3}{4} \cdot \frac{4+1}{10+7} \cdot \frac{1}{10+7} = \frac{1}{Z} \cdot \underline{0.013}$$

$$P(-|\text{flu influenza}) = \frac{1}{Z} P(-) P(\text{flu}|-) P(\text{influenza}|-) = \frac{1}{Z} \cdot \frac{1}{4} \cdot \frac{1+1}{5+7} \cdot \frac{1}{5+7} = \frac{1}{Z} \cdot \underline{0.003}$$

where $Z = P(\text{flu influenza})$, which can be ignored (why?).

The predicted label is the one with the highest probability:

$$\arg \max_{y \in \{+, -\}} P(y|\text{flu influenza}) = +$$

Laplace smoothing again

The extra smoothing words are sometimes called *pseudowords*.

Pseudowords don't have to be integers; one might e.g. add 0.01 pseudowords per class.

In our example:

$$P(+)P(\text{flu}|+)P(\text{influenza}|+) = \frac{3}{4} \cdot \frac{4+0.01}{10+0.07} \cdot \frac{0.01}{10+0.07} = \underline{0.0003}$$

$$P(-)P(\text{flu}|-)P(\text{influenza}|-) = \frac{1}{4} \cdot \frac{1+0.01}{5+0.07} \cdot \frac{0.01}{5+0.07} = \underline{0.00001}$$

A practicality

At test time, one might encounter words that do not appear in the training data.

(Many ML packages do not allow this.)

Recommended approach: Ignore those unknown words.

will be equally unknown for each class

DD2418 Language Engineering

5c: Evaluation of text classification

Johan Boye, KTH

Evaluation

Evaluation involves testing the hypothesis h on a development set $\{(x_1, y_1), \dots, (x_k, y_k)\}$.

Confusion matrix:

		Real label y	
		<i>spam</i>	<i>ham</i>
$h(x)$	<i>spam</i>	102	48
	<i>ham</i>	27	323

Positive examples (for *spam*): $102+48 = 150$

Negative examples: $27+323 = 350$

True positives: 102

False positives: 48

True negatives: 323

False negatives: 27

Evaluation metrics

Accuracy (binary case):

$$\frac{TP + TN}{TP + TN + FP + FN}$$

Accuracy (generally):

$$\frac{\text{correctly classified examples}}{\text{all examples}}$$

Precision:

$$\frac{TP}{TP + FP}$$

Recall:

$$\frac{TP}{TP + FN}$$

Evaluation quiz

Evaluation involves testing the hypothesis h on a development set $\{(x_1, y_1), \dots, (x_k, y_k)\}$.

Confusion matrix:

		Real label y	
		<i>spam</i>	<i>ham</i>
$h(x)$	<i>spam</i>	102	48
	<i>ham</i>	27	323

Compute accuracy of the classifier, as well as precision and recall (for the spam class).

Evaluation quiz

Evaluation involves testing the hypothesis h on a development set $\{(x_1, y_1), \dots, (x_k, y_k)\}$.

Confusion matrix:

		Real label y	
		<i>spam</i>	<i>ham</i>
$h(x)$	<i>spam</i>	102	48
	<i>ham</i>	27	323

Accuracy: $(102+323) / (102+323+48+27) = .85$

Precision: $102/(102+48) = .68$

Recall: $102/(102+27) = .79$

Precision-recall trade-off

What is most important, precision or recall?

(Arithmetic) average precision/recall:

$$\frac{P + R}{2}$$

More useful: Geometric mean (F-score):

$$\frac{1}{\frac{1}{2}(\frac{1}{P} + \frac{1}{R})} = \frac{2PR}{P + R}$$

Precision-recall trade-off

(Arithmetic) average precision/recall:

$$\frac{P + R}{2}$$

If $P = R = 0.5$, then $\frac{P+R}{2} = 0.5$.

If $P = 0.01$ and $R = 0.99$, then $\frac{P+R}{2} = 0.5$.

More useful: Geometric mean (F-score):

$$\frac{1}{\frac{1}{2}(\frac{1}{P} + \frac{1}{R})} = \frac{2PR}{P + R}$$

If $P = R = 0.5$, then $\frac{2PR}{P+R} = 0.5$.

If $P = 0.01$ and $R = 0.99$, then $\frac{2PR}{P+R} = 0.0198$.

Non-binary case

Classify news as *Sports*, *Politics*, or *Economy*.

x	$h(x)$	y
1	S	S
2	P	S
3	E	S
4	S	S
5	P	P
6	E	P
7	S	P
8	P	P
9	E	E
10	S	E

Confusion matrix

		Real label y		
		S	P	E
$h(x)$	S	2	1	1
	P	1	2	0
	E	1	1	1

Non-binary case, quiz

Given the confusion matrix on the previous slide, compute these quantities:

Accuracy _____

Precision, class S _____

Precision, class P _____

Precision, class E _____

Recall, class S _____

Recall, class P _____

Recall, class E _____

Non-binary case, quiz

Given the confusion matrix on the previous slide, compute these quantities:

Accuracy	0.5
Precision, class S	0.5
Precision, class P	0.67
Precision, class E	0.33
Recall, class S	0.5
Recall, class P	0.5
Recall, class E	0.5

Evaluation methodology

3 sets: *training*, *development* and *test* sets.

The three sets should all be distinct!

- Otherwise the model will be *over-trained*.
- Results will be unrealistically good.
- Model will perform poorly on new data.

First, set aside 10% of the data as the final test set.

validation set

Then train on 80% of the data, evaluate on 10%. Iterate.

To minimize random effects, *n-fold cross-validation* can be used.

Finally, evaluate on the test set. This is **the** result.

Hyperparameters

Typically, a machine learning method has a number of hyperparameters that influence the result.

- e.g. the number of layers in a neural network, and the number of neurons in each layer

Typical choices in a language engineering setting:

- Case folding, yes/no?
- Lemmatization, yes/no?
- Remove certain words (stop words)?
- Context window size?
- Normalization of numbers, years, dates, ...?
- Tokenization, what is a word?
- ...

How good is 80% accuracy?

If your data contains 90% ham and 10% spam, then a classifier returning “*ham*” all the time has 90% accuracy.

There might be other obvious base lines (e.g. check for the term “*weight loss*”), which beats the random baseline.

Make sure that you compare your results to baseline results!

- ... otherwise you risk fooling yourself!

Gold standard

The *gold standard* is the correct labeling according to a human expert.

Very useful, but can be hard/expensive to obtain.